

Prelim 2

Computer Science 3410, Cornell University
Spring 2016, Prof. Bracy

5 May 2016

Instructions:

- Please turn off and stow away all electronic devices.
- This is a **closed book and notes** examination. You **may** use the 3-sided reference given to you by the instructor.
- You have **120 minutes** to complete 121 points. Use your time accordingly.
- There are 9 problems and 14 pages. Make sure you have the whole exam.
- To receive partial credit you must show your work. If you believe a question is open to interpretation, then please ask us about it! Please state any assumptions you make.

Good luck!

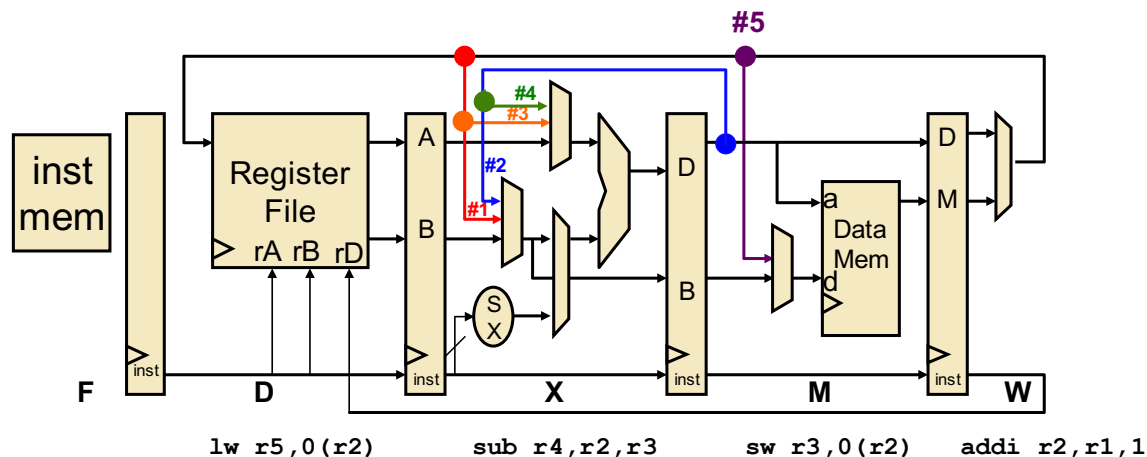
Problem	Points	Score
1	0	
2	16	
3	8	
4	16	
5	10	
6	24	
7	15	
8	8	
9	24	
Total	121	

Your Name _____

1. Write your name here and NetID at the top of each odd page or face -2 pt penalty. [0 pts]

2. Data and Control Hazards [16 pts]

- (a) [8 pts] **Forwarding.** In the pipeline below, the register file is read in the 1st half of the cycle and written in the 2nd half of the cycle. Bypasses #1-#5 are shaded. In the D/X latch, the 1st operand to an ALU instruction is stored in the A slot of the latch. The 2nd operand is stored in the B slot of the latch.



For each input specified below, state where the instruction gets its data. 1-5 refer to the bypasses shown in the above figure. RF stands for register file. ERROR means that the instruction cannot possibly get its data in order to execute correctly given the processor above and the timing specified in the diagram.

	Circle One						
r2 of the sw	#1	#2	#3	#4	#5	RF	ERROR
r2 of the sub	#1	#2	#3	#4	#5	RF	ERROR
r3 of the sub	#1	#2	#3	#4	#5	RF	ERROR
r2 of the lw	#1	#2	#3	#4	#5	RF	ERROR

- (b) [8 pts] **Control Dependences.** Suppose we had asked you to implement a 2-instruction control delay slot for Project 2. In other words, the next *two* instructions immediately following a control instruction will always be executed. Which of the following statements about this concept are true? Check all that apply.

- ☐ 2 delay slots would be harder to implement in Logisim than 1.
☐ It will be harder for a compiler to fill the 2nd delay slot than to fill the 1st.
☐ Processors with 1 delay slot benefit *more* from branch prediction than one with 2 delay slots.
☐ A processor with 2 delay slots might support a faster clock than one with only 1.

3. ISAs [8 pts]

For each multiple choice question there is only *one* correct answer.

(a) [2 pts] What is one advantage of a CISC ISA?

- A It naturally supports a faster clock.
- B Instructions are easier to decode.
- C The static footprint of the code will be smaller.
- D The code is easier for a compiler to optimize.
- E You have a lot of registers to use.

Your Answer: _____

(b) [2 pts] What processor feature does a compiler **not** need to know about?

- A There is a control delay slot.
- B The number of inputs each instruction can have.
- C The processor performs statically-scheduled multiple issue execution.
- D The processor performs dynamically-scheduled multiple issue execution.
- E Whether the processor supports predication.

Your Answer: _____

(c) [2 pts] What is **not** a good reason to limit the number of instructions offered by a particular ISA?

- A Offering more instructions leads to executables with more static instructions.
- B More complicated instructions are difficult to pipeline.
- C More instructions may lead to needing more bits in the opcode field which could make static instructions larger.
- D It is very difficult to ever stop supporting a particular instruction once it has been added to the ISA.
- E More instructions will complicate the decode logic of the processor.

Your Answer: _____

(d) [2 pts] Which of the following statements about ISAs is **true**?

- A Because clock frequencies are no longer increasing, a return to CISC ISAs makes sense.
- B Which cache coherence protocol is implemented needs to be specified in the ISA.
- C Atomic operations need to be specified in the ISA.
- D ISA design has stagnated; no new ISAs of any commercial significance have appeared in the past 10 years.
- E A system call is OS-specific and not an ISA-level feature.

Your Answer: _____

4. Calling Conventions [16 pts]

The following assembly is the body of a function compiled for a MIPS processor with a control delay slot:

```
Body:    addiu $s3, $a0, 0
          addiu $s5, $a1, 0
          addiu $s6, $a2, 0
Loop:    sll $t1,$s3,2
          add $t1,$t1,$s6
          lw  $t0,0($t1)
          bne $t0,$s5, Exit
          nop
          addi $s3,$s3,1
          j  Loop
          nop
Exit:    addiu $a0, $s3, 0
          jal record
          nop
          addiu $v0, $s3, 0
```

Write the prologue and epilogue for this body. Use the class calling conventions.

PROLOGUE:

EPILOGUE:

5. Linkers and Loaders. [10 pts]

Below is a modified version of `heaptest.c` from Project 4:

```
-----  
#include <stdio.h>  
#include heaplib.h  
  
#define HEAP_SIZE 16  
static int ARR_SIZE = 4;  
  
int main() {  
  
    char heap[HEAP_SIZE];  
    hl_init(heap, HEAP_SIZE * sizeof(char));  
    char* ptr = (char *) hl_alloc(heap, ARR_SIZE * sizeof(char));  
  
    ptr[0] = 'h';  
    ptr[1] = 'i';  
    ptr[2] = '\\0';  
    printf("%s\\n", ptr);  
    return 0;  
}  
-----
```

Where does the assembler place each of the following symbols from the above program in the object file that it creates?

Choose one:

- (a) Text Segment
- (b) Data Segment
- (c) Exported reference in the symbol table
- (d) Imported reference in the symbol table
- (e) None of the above

Circle One:

- | | | | | | |
|---------------|---|---|---|---|---|
| (1) HEAP_SIZE | A | B | C | D | E |
| (2) ARR_SIZE | A | B | C | D | E |
| (3) heap[] | A | B | C | D | E |
| (4) hl_init | A | B | C | D | E |
| (5) 'h' | A | B | C | D | E |

6. Caches [24 pts] (parts a–d)

- (a) [9 pts] For each workload, choose the best block size for your cache among the choices given. Assume that integers and pointers are all 4 bytes each. Use the intuitions you have developed in this class to decide what **best** means for a cache.

Choose one: (a) 1 byte (b) 4 bytes (c) 8 bytes (d) 16 bytes (e) 32 bytes

BEST BLOCK SIZE (Circle One)

```
int scores[NUM_STUDENTS] = 0;          A      B      C      D      E
```

```
int sum = 0;
for (i = 0; i < NUM_STUDENTS; i++) {
    sum += scores[i];
}
```

```
typedef struct _item_t {
    int value;
    item_t *next;
    char *name;
} item_t;                                A      B      C      D      E
```

```
int sum = 0;
item_t *curr = list_head;
while (curr != NULL) {
    sum += curr->value;
    curr = curr->next;
}
```

```
// H = 16, W = 16                                Assume a 256 byte cache.
```

```
int A[H][W];
```

```
for(x=0; x < W; x++)
    for(y=0; y < H; y++)
        sum += A[y][x];                    A      B      C      D      E
```

- (b) [9 pts] **Memory Access Time.** Processor A is a 1 GHz processor (1 cycle = 1 ns). There is an L1 cache with a 1 cycle access time and a 50% hit rate. There is an L2 cache with a 10 cycle access time and a 90% hit rate. It takes 100 ns to access memory.

(3 points) What is the average memory access time of Processor A in ns?

Processor B attempts to improve upon the hit rate of Processor A by making the L1 cache larger. The new hit rate is 90%. In order to maintain a 1 cycle access time, Processor B can only run at 500 MHz (1 cycle = 2 ns).

(3 points) What is the average memory access time of Processor B in ns?

(3 points) Which processor would you buy, and why?

-
- (c) [4 pts] Consider a 32-bit processor without virtual memory. It has an 24K, 3-way set associative cache with a 32-byte line size.

How many index bits are there? Your Answer: _____

How many tag bits are there? Your Answer: _____

- (d) [2 pts] **Cache Design.** Consider a 32-bit processor without virtual memory; its cache has 1 byte cache lines. What would happen if the most significant bits of the address were used as the cache index and the least significant bits of the address were used as the cache tag? What is the best assessment of the result of this change?
- A The correctness of the processor will be compromised.
 - B The performance of the processor will be compromised.
 - C The correctness and the performance of the processor will be compromised.
 - D Neither the correctness nor the performance will be compromised.

Your Answer: _____

7. TLB Performance Optimizations [15 pts]

On machines with virtual memory, before an instruction can be fetched, the PC must be translated from a virtual address to a physical address. This means the TLB and the I-Cache must be accessed in serial (TLB first, I-Cache second). One performance optimization employed by many processors is to access the TLB and the I-Cache **in parallel**. The cache index is extracted from the *virtual address* and is used to perform the cache lookup. At the same time, the Virtual Page Number (VPN) is extracted from the virtual address and is used to perform the TLB lookup. After the Physical Page Number (PPN) is read from the TLB, it is used to construct the tag which is then used to determine whether there is a tag match in the cache (*i.e.*, a cache hit).


Consider a machine with 16-bit virtual and physical addresses. For cache lookups, the bottom 3 bits of the physical address are used as the byte offset. The next 4 bits of the physical address are the index, and the remaining 9 bits of the physical address are the tag. Suppose pages are 256 bytes each.

- (a) [2 pts] How many bits are needed to represent the page offset?

Your Answer: _____

- (b) [2 pts] The cache is physically addressed. Why is it okay to use the index bits from the virtual address to perform the cache lookup? Be brief.

You may find the following diagrams useful, but they will not be graded:



(cont'd from previous page) Consider the virtual address **1000 1000 1000 1000**.

- (c) [2 pts] What is the value of the VPN in binary?

Your Answer: _____

- (d) [3 pts] Suppose that this VPN's PPN (which is stored in the TLB) is all 0s. (For example, if the PPN were 4 bits long, it would be 0000.) What tag value (*in binary*) should be used for the tag match in the cache?


Your Answer: _____

- (e) [6 pts] You want to make the cache 4x larger. What strategy would prevent the parallel access of the TLB and I-Cache? *Answer all that apply.*

- A Increasing the associativity of the cache by 4x.
- B Increasing the line size of the cache by 4x.
- C Increasing the number of entries in the cache by 4x.

Your Answer: _____

You may find the following diagrams useful, but they will not be graded:



8. Exceptional Control Flow [8 pts] (parts a–d)

(a) [2 pts] What does *asynchronous* mean in the context of exceptional control flow?

Below are 4 simplified descriptions of how different exceptional situations are handled:

Sequence A:

- current instruction (at PC) triggers handler
- control passes to handler
- execution never returns to the user process

Sequence B:

- current instruction (at PC) triggers handler
- control passes to handler
- execution returns to user process
- current instruction (at PC) executes once more

Sequence C:

- current instruction (at PC) completes
- control passes to handler
- execution returns to user process
- next instruction (at PC+4) executes

Sequence D:

- current instruction (at PC) triggers handler
- control passes to handler
- execution returns to user process
- next instruction (at PC+4) executes

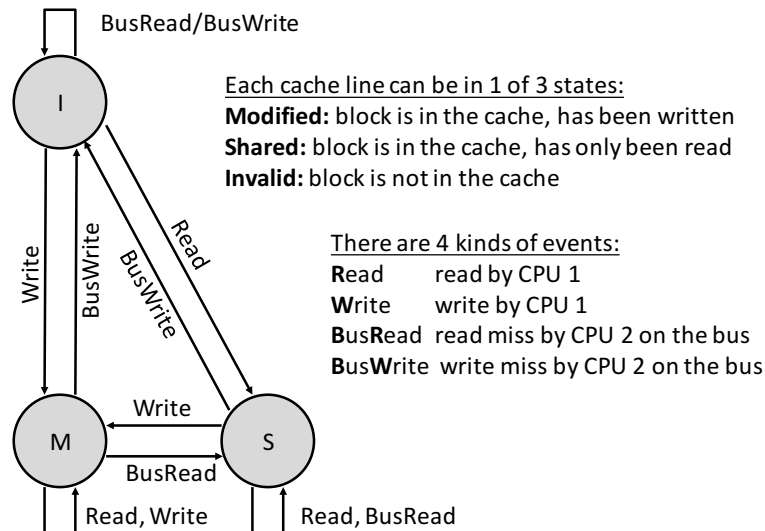
Circle One:

- | | | | | |
|--|---|---|---|---|
| (b) [2 pts] Which sequence best describes a system call ? | A | B | C | D |
| (c) [2 pts] Which sequence best describes a fault ? | A | B | C | D |
| (d) [2 pts] Which sequence best describes an interrupt ? | A | B | C | D |

9. Multicore [24 pts] (parts a-c)

- (a) [12 pts] This question concerns a direct mapped 32-byte cache with 8-byte blocks. This is in a multiprocessor. You are simulating the cache associated with CPU 0. CPU 0 shares a bus with CPU 1. Reference addresses come from both the local processor (CPU 0) and other processors (CPU 1 via the bus). The cache uses an LRU replacement policy but would *always* replace an invalid block before a valid one.

This multiprocessor uses a simplified 3-state MSI protocol as defined here:



There are 5 kinds of misses: cold, conflict, capacity, upgrade and coherence.

The initial contents of the cache (base block address : state) are given in the first line of the table. Assume that the initial blocks were accessed in the following order: 000, 050, 020, 070. For each Event, fill in:

- (1) whether the event results in a Hit or Miss if applicable. (If a miss, what kind?)
- (2) what bus event (if any) must be generated for CPU 1 to see

If the Event triggers a change in the cache, show that changed state on the next line. To simplify the notation, we use an *octal* block address. For example: $046_{octal} = 000\ 100\ 110_{binary}$. The first access has been done for you.

Cache state prior to access						Event addr:action	Hit or Miss? (if applicable)	Bus Event addr:action (if applicable)
Set 0			Set 1					
Way 0	LRU	Way 1	Way 0	LRU	Way 1			
000:S	0	020:S	050:M	0	070:S	051:R	hit	—
				1		034:W		
						024:W		
						020:BW		
						072:R		
						022:W		

- (b) [4 pts] Most modern cache coherence protocols also include an E (Exclusive) state. A cache line in the Exclusive state means (1) this local cache has the *only* cached copy of the data and (2) this copy is clean. Under the MESI protocol, cache lines begin in the Invalid state and a Read can transition from I to E or from I to S.

(2 points) Under what conditions would a Read transition the cache line from the I state to the S state? (Limit your answer to 1 sentence.)

(2 points) What is the E state designed to prevent? (Why is it useful?) Do not write more than 2 sentences.

- (c) [8 pts] Many CPU architectures provide a Compare-And-Swap (CAS) instruction. It is a single, atomic assembly instruction that can be called with three register inputs (one of which is also a register output) with the semantics:

CAS *addr*, *compare_value*, *new_value*

- *addr*: an address in memory
- *compare_value*: as an input, it is the value you expect to be at *addr*; as an output, it is 1 for success (meaning the compare was a match and the new-value was stored at *addr*) and 0 for failure (meaning the compare was not a match and no new value was written to *addr*)
- *new_value*: the new value you will write to *addr* *if the compare_value matched*

Or if you prefer C to English:

```
ATOMIC void CAS(int *addr, int *comp_val, int new_val) {
    if (*addr != *comp_val) {
        *comp_val = 0;
    } else {
        *addr = new_val;
        *comp_val = 1;
    }
}
```

Below is the MIPS implementation of `mutex_lock`. Rewrite the body of this function so that it uses a CAS instruction instead of LL/SC. The function behavior should not change. The address of the lock is the argument to the function.

```
mutex_lock(int *m) {
    test: LI $t0, 1
          LL $t1, 0($a0)
          BNEZ $t1, test
          SC $t0, 0($a0)
          BEQZ $t0, test
}
```

```
mutex_lock(int *m) {
```

```
}
```