

- Write your NetID at the top of each physical page of the exam **or face -2 pt penalty**.
- Please turn off and stow away all electronic devices. You may not use them for any reason for the entirety of the exam. Do not bring them with you if you leave the room temporarily.
- This is a closed book and notes examination. You may use the 3-sided reference provided.
- To receive partial credit you must show your work. If you believe a question is open to interpretation, then please ask us about it! Please state any assumptions you make.
- There are 8 problems and 12 pages. Make sure you have the whole exam. You may use the back-side of each page as scratch paper, but it will *not* be graded.
- You have **90 minutes** to complete 90 points. Use your time accordingly.

Problem	Topic	Points	Score
1	Multiple Choice	18	
2	Translation	5	
3	Think of the Possibilities!	8	
4	The Need for Speed	12	
5	Call Me Maybe	15	
6	Batman and Robin	12	
7	Cache on Hand?	10	
8	Miss Congeniality	10	
Total		90	

Your Name \_\_\_\_\_

Your NetID \_\_\_\_\_

1. Multiple Choice [18 pts] There is only 1 correct answer per question. If you write down multiple answers, we will only grade the first one.

(a) [2 pts] **Multi Level Page Tables.** What is *not* a reason for implementing a multi-level page table?

- A. Multi-Level Page Tables support a faster translation than Single Level Page Tables.
- B. Unlike Single Level Page Tables, Multi-Level Page tables do not need to be contiguous in memory.
- C. Unlike Single Level Page Tables, Multi-Level Page tables do not need to have a PTE allocated for every virtual page number.
- D. Unlike Single Level Page Tables, Multi-Level Page tables allow two processes to map distinct virtual page numbers to the same physical page.
- E. In the common case, a Multi-Level Page Tables takes up fewer pages in memory than Single Level Page Tables.

Your Answer: \_\_\_\_\_

(b) [2 pts] **Conflict Misses.** You can have a conflict miss in a fully associative cache.

- A. True
- B. False
- C. Cannot be answered with the information given

Your Answer: \_\_\_\_\_

(c) [2 pts] **Coherence.** Two cores sharing a single cache still need a coherence protocol for their shared cache in order to maintain coherence.

- A. True
- B. False
- C. Cannot be answered with the information given

Your Answer: \_\_\_\_\_

(d) [2 pts] **C Basics.** What value is printed in the last line of `main`?

```
void set1(int val1, int *val2) {
    *val2 = val1 + 6;
}

int set2(int val1, int *val2) {
    return (*val2 + 3);
}
```

- A. 8
- B. 11
- C. 14
- D. It depends on the address of x.
- E. None of the above.

Your Answer: \_\_\_\_\_

```
int main() {
    int x = 5;
    int *y = &x;
    set1(set2(x, y), y);
    printf("%d\n", x);
}
```

(e) [2 pts] **VI Protocol.** The VI protocol's biggest drawback is the upgrade misses.

- A. True
- B. False
- C. Cannot be answered with the information given

Your Answer: \_\_\_\_\_

(f) [2 pts] **Asynchronicity.** Which of these exceptional events are asynchronous?

- A. Software Exception
- B. Hardware Interrupt
- C. System Call
- D. A and C
- E. B and C

Your Answer: \_\_\_\_\_

(g) [2 pts]     **Synchronization.** Which of the following statements about Synchronization is *false*?

- A. LL/SC is helpful when cooperating threads of a parallel program need to synchronize to get proper behavior for reading and writing shared data.
- B. LL/SC is helpful When cooperating processes on a uniprocessor need to synchronize for reading and writing shared data.
- C. LL/SC are an example of hardware support for synchronization.
- D. A programmer can benefit from LL/SC instructions without understanding them well enough to use them correctly at the assembly level.
- E. LL/SC are the instructions a programmer should use to start and end a critical section.

Your Answer: \_\_\_\_\_

(h) [2 pts]     **Exceptions.** A five-stage pipeline (IF, ID, EX, MEM, WB) executes the following instruction sequence:

```
A    lw $1, 0($2)    # page fault
B    XXX $1, $2, $1  # undefined instruction
C    add $1, $2, $1  # arithmetic overflow
D    sub $1, $2, $1  # hardware error
```

Which exception (A, B, C, or D) should be recognized first in the above sequence?

Your Answer: \_\_\_\_\_

(i) [2 pts]     **Saving State.** In order to preserve the processor state of the user process, the OS handler saves both caller- and callee- saved registers after which of the following?

- A. Software Exception
- B. Hardware Interrupt
- C. System Call
- D. A and B
- E. A, B, and C

Your Answer: \_\_\_\_\_

2. Translation [5 pts]

Consider a machine with a 20-bit virtual address, a 16-bit physical address and 2KB pages in memory. Each page table entry is 4 bytes.

- (a) [1 pt]     How large is the page offset?
  
  
  
  
  
  
  
  
  
  
- (b) [2 pts]     How many pages in memory would be required to store the entire single-level page table?
  
  
  
  
  
  
  
  
  
  
- (c) [2 pts]     What would the benefit be of having a 2 level page table instead?  
Be specific.

3. Think of the Possibilities! [8 pts]

For each combination, fill in the table indicating whether the combination is:

- A. Possible
- B. Impossible
- C. Possible but never detected
- D. Cannot say with the given information

If you believe it is possible, state under what circumstance this would happen.  
If it is not possible, not detectable, or not knowable, explain why.

<b>TLB</b>	<b>Page Table</b>	<b>Cache</b>	Possible? (A-D)	<i>Explanation</i>
Miss	Hit	Hit		
Hit	Hit	Miss		
Miss	Miss	Miss		
Hit	Miss	Miss		

## 4. The Need for Speed [12 pts] (parts a–d)

Consider a 1 GHz (1 cycle = 1 ns) pipelined MIPS processor with an instruction mix of 20% loads and 80% integer operations. The processor has an instruction cache with a 90% hit rate and a data cache with an 80% hit rate. Both caches are accessed in 1 cycle. It takes 10ns to access DRAM. Integer operations take 1 cycle to execute. In optimal circumstances, an instruction finishes every cycle, yielding a CPI of 1.

- (a) [3 pts] What is the CPI of this machine?
- (b) [3 pts] **Modification 1:** You double the size of the data cache, which increases its hit rate to 90%. In order to maintain a 1 cycle access time, the clock rate is slowed to 500MHz. DRAM speed is unchanged. What is the CPI of this modified machine?
- (c) [4 pts] **Modification 2:** Instead of modifying the data cache (as in the previous section), you add a *unified* level 2 cache with a 4 cycle access time and a 90% hit rate. This L2 cache will be accessed *in parallel* with the first level caches. (L1 hits do not wait for L2 results.) Clock frequency remains unchanged from the original processor. What is the CPI of this modified machine?
- (d) [2 pts] Which processor is faster: the original, Modified 1, or Modified 2? Explain.

5. Call Me Maybe [15 pts] Consider the following typedef and function in C. The `main` calls function `get_point` which calls function `scale_y`, using the 3410 Calling Conventions. Note that the dot (`.`) operator takes higher precedence than the ampersand (`&`) operator.

```
typedef struct point_t {      1 void get_point(int scale, int val) {
    int x;                    2     point_t p1, p2;
    int y;                    3     p1.y = 1 + val;
} rect_t;                     4     p2.y = 3;
                                5     scale_y(scale, &p1.y, val);
                                6     p1.x = scale * p1.x * p2.y;
                                7     return &p1;
                                8 }
```

- (a) [12 pts] In as few lines of code as possible, write the assembly for `get_point`, beginning with the prologue and **ending with line 5**. Do **not** use any `s` registers. To be eligible for partial credit, a line must have a comment, explaining what the line accomplishes.



- (b) [1 pt] Which line in `get_point` is bad C code? \_\_\_\_\_
- (c) [2 pts] Explain in 1-2 sentences why.  
(*We will only read your first 2 sentences.*)

6. Batman and Robin [12 pts] You maybe have known the word *dynamic* from expressions like 'the dynamic duo' but in CS 3410 you have learned some very specific, technical meanings of the word dynamic. Choose **THREE** of the following 5 instances of the word **dynamic**. Explain what is meant by **dynamic** (as opposed to static) and answer the follow-up question. If you answer more than 3, only the first 3 will be graded.

- **Dynamic Random Access Memory (DRAM).**

- What does dynamic mean?

- What is one way in which DRAM is better than SRAM?

- **Dynamic Instruction Count.**

- What does dynamic mean?

- Why might one care more about the *static* instruction count?

- **Dynamically Linking.**

- What does dynamic mean?

- What is one advantage of a dynamic linking?

- **Dynamic Memory Allocation.**

- What does dynamic mean?

- What is one advantage of allocating memory dynamically?

- **Dynamic Scheduling.**

- What does dynamic mean?

- What is one advantage of a static scheduling?

# 7. Cache on Hand? [10 pts]

The following problem concerns basic cache lookups.

- The memory is byte addressable.
- Memory accesses are to **1-byte words** (not 4-byte words).
- Physical addresses are 12 bits wide.
- The cache is 4-way set associative, with a 2-byte block size and 32 total lines.

- (a) [3 pts] The box below shows the format of a physical address. Indicate how each bit would be used by filling in the boxes with ***O** for block offset*, ***I** for cache index*, and ***T** for cache tag*).

11	10	9	8	7	6	5	4	3	2	1	0

In this table, **all numbers are hexadecimal**. The contents of the cache are as follows:

4-way Set Associative Cache																
Index	Tag	Valid	Byte 0	Byte 1	Tag	Valid	Byte 0	Byte 1	Tag	Valid	Byte 0	Byte 1	Tag	Valid	Byte 0	Byte 1
0	29	0	34	29	87	0	39	AE	7D	1	68	F2	8B	1	64	38
1	F3	1	0D	8F	3D	1	0C	3A	4A	1	A4	DB	D9	1	A5	3C
2	A7	1	E2	04	AB	1	D2	04	E3	0	3C	A4	01	0	EE	05
3	3B	0	AC	1F	E0	0	B5	70	3B	1	66	95	37	1	49	F3
4	80	1	60	35	2B	0	19	57	49	1	8D	0E	00	0	70	AB
5	EA	1	B4	17	CC	1	67	DB	8A	0	DE	AA	18	1	2C	D3
6	1C	0	3F	A4	01	0	3A	C1	F0	0	20	13	7F	1	DF	05
7	0F	0	00	FF	AF	1	B1	5F	99	0	AC	96	3A	1	22	79

For the given physical address, indicate the cache entry accessed and the cache byte value returned **in hex**. Indicate whether a cache miss occurs.

If there is a cache miss, enter “-” for “Cache Byte returned”.

**Physical address: 3B6**

- (b) [2 pts] Physical address format (one bit per box)

11	10	9	8	7	6	5	4	3	2	1	0

- (c) [5 pts] Physical memory reference

Parameter	Value
Cache Offset (CO)	0x
Cache Index (CI)	0x
Cache Tag (CT)	0x
Cache Hit? (Y/N)	
Cache Byte returned	0x

### 8. Miss Congeniality [10 pts]

Consider a dual core processor with 4-bit addresses. Each core has a private **8B** cache with 2 byte cache lines. Core 0's cache is fully associative. Core 1's cache is direct mapped.

The caches use an MSI coherence protocol. For simplicity, the diagram below shows only the full 4-bit address of the block-aligned memory address. Both caches happen to have been filled from left to right (Core 0's cache had accesses in the order: 1011, 1111, 0010, 0000. Core 1's cache had accesses in the order: 0000, 1010, 0101, 1111. The diagram below shows memory references at time 100, 101, and 102. Each reference has format X:ABCD:Y where X indicates the Core that the reference originated from and Y indicates whether it is a read or a write.

CORE 0 Cache state prior to access				t	Memory Reference	CORE 1 Cache state prior to access			
Way 0	Way 1	Way 2	Way 3			Set 0	Set 1	Set 2	Set 3
1010:S	1110:S	0010:S	0000:S	100	0:0011:R	0000:S	1010:S	0100:M	1110:S
0010:S				101	1:1101:R			1100:S	
				102	1:1111:W				
	1110:I			103	???				1110:M

Your job is to suggest 6 different Memory References that at time 103 that would be categorized in each of the following ways. The first one has been done for you. If no reference receives the desired categorization, please write 'not possible'. Note that your references do *not* happen one after the other. They represent 6 possible references for time 103 only.

Goal	t=103 Memory Reference
Hit	1:1110:R
Cold Miss	
Conflict Miss	
Capacity Miss	
Coherence Miss	
Upgrade Miss	